

Predicting Taxi Passenger Generosity: A Machine Learning Model for Tip Behavior

The main dataset is from New York City Taxi & Limousine Commission. The goal is to build a model to predict whether a taxi customer is a generous tipper.

1. Packages and Libraries

```
In [4]: import numpy as np
import pandas as pd

import matplotlib.pyplot as plt

from sklearn.model_selection import GridSearchCV, train_test_split
from sklearn.metrics import roc_auc_score, roc_curve
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix,
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from xgboost import plot_importance

import pickle

pd.options.mode.chained_assignment = None
```

2. Datasets

```
In [5]: # Load datasets
pd.set_option('display.max_columns', None)
df0 = pd.read_csv('2017_Taxi.csv')
nyc_preds_means = pd.read_csv('predicted_means.csv')
```

```
In [6]: df0.head(10)
```

Out[6]:

	Unnamed: 0	VendorID	tpep_pickup_datetime	tpep_dropoff_datetime	passenger_count	trip_distance	Ratecode
0	24870114	2	03/25/2017 8:55:43 AM	03/25/2017 9:09:47 AM	6	3.34	
1	35634249	1	04/11/2017 2:53:28 PM	04/11/2017 3:19:58 PM	1	1.80	
2	106203690	1	12/15/2017 7:26:56 AM	12/15/2017 7:34:08 AM	1	1.00	
3	38942136	2	05/07/2017 1:17:59 PM	05/07/2017 1:48:14 PM	1	3.70	
4	30841670	2	04/15/2017 11:32:20 PM	04/15/2017 11:49:03 PM	1	4.37	
5	23345809	2	03/25/2017 8:34:11 PM	03/25/2017 8:42:11 PM	6	2.30	
6	37660487	2	05/03/2017 7:04:09 PM	05/03/2017 8:03:47 PM	1	12.83	
7	69059411	2	08/15/2017 5:41:06 PM	08/15/2017 6:03:05 PM	1	2.98	
8	8433159	2	02/04/2017 4:17:07 PM	02/04/2017 4:29:14 PM	1	1.20	
9	95294817	1	11/10/2017 3:20:29 PM	11/10/2017 3:40:55 PM	1	1.60	

◀ ▶

In [7]: `nyc_preds_means.head(10)`

Out[7]:

	mean_duration	mean_distance	predicted_fare
0	22.847222	3.521667	16.434245
1	24.470370	3.108889	16.052218
2	7.250000	0.881429	7.053706
3	30.250000	3.700000	18.731650
4	14.616667	4.435000	15.845642
5	11.855376	2.052258	10.441351
6	59.633333	12.830000	45.374542
7	26.437500	4.022500	18.555128
8	7.873457	1.019259	7.151511
9	10.541111	1.580000	9.122755

In [8]:

```
# Merge datasets
df0 = df0.merge(nyc_preds_means,
                 left_index=True,
                 right_index=True)
df0.head()
```

Out[8]:

	Unnamed: 0	VendorID	tpep_pickup_datetime	tpep_dropoff_datetime	passenger_count	trip_distance	Ratecode
0	24870114	2	03/25/2017 8:55:43 AM	03/25/2017 9:09:47 AM	6	3.34	
1	35634249	1	04/11/2017 2:53:28 PM	04/11/2017 3:19:58 PM	1	1.80	
2	106203690	1	12/15/2017 7:26:56 AM	12/15/2017 7:34:08 AM	1	1.00	
3	38942136	2	05/07/2017 1:17:59 PM	05/07/2017 1:48:14 PM	1	3.70	
4	30841670	2	04/15/2017 11:32:20 PM	04/15/2017 11:49:03 PM	1	4.37	

In [9]:

```
df0.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 22699 entries, 0 to 22698
Data columns (total 21 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Unnamed: 0        22699 non-null   int64  
 1   VendorID         22699 non-null   int64  
 2   tpep_pickup_datetime  22699 non-null   object 
 3   tpep_dropoff_datetime  22699 non-null   object 
 4   passenger_count    22699 non-null   int64  
 5   trip_distance      22699 non-null   float64 
 6   RatecodeID         22699 non-null   int64  
 7   store_and_fwd_flag  22699 non-null   object 
 8   PULocationID       22699 non-null   int64  
 9   DOLocationID       22699 non-null   int64  
 10  payment_type       22699 non-null   int64  
 11  fare_amount        22699 non-null   float64 
 12  extra              22699 non-null   float64 
 13  mta_tax             22699 non-null   float64 
 14  tip_amount          22699 non-null   float64 
 15  tolls_amount        22699 non-null   float64 
 16  improvement_surcharge  22699 non-null   float64 
 17  total_amount         22699 non-null   float64 
 18  mean_duration        22699 non-null   float64 
 19  mean_distance         22699 non-null   float64 
 20  predicted_fare       22699 non-null   float64 

dtypes: float64(11), int64(7), object(3)
memory usage: 3.6+ MB
```

3. Feature Engineering

Target Variable Y

In [10]:

```
# Customers who pay cash generally have a tip amount of $0. Filtering for only the customers using credit cards
df1 = df0[df0['payment_type']==1]
```

In [11]:

```
# Calculating tip percent, Rounding for floating-point arithmetic
df1['tip_percent'] = round(df1['tip_amount'] / (df1['total_amount'] - df1['tip_amount']), 4)
df1['tip_percent']
```

```
Out[11]: 0      0.2000
         1      0.2381
         2      0.1986
         3      0.3000
         5      0.2000
         ..
        22692    0.1995
        22693    0.2000
        22695    0.2500
        22697    0.1504
        22698    0.1992
Name: tip_percent, Length: 15265, dtype: float64
```

```
In [12]: # Creating target variable with customer tipping >= 20% as "generous", encoding categorical variable.
df1['generous'] = (df1['tip_percent'] >= 0.20).astype(int)
df1['generous']
```

```
Out[12]: 0      1
         1      1
         2      0
         3      1
         5      1
         ..
        22692    0
        22693    1
        22695    1
        22697    0
        22698    0
Name: generous, Length: 15265, dtype: int64
```

Datetime Management

```
In [13]: df1['tpep_pickup_datetime'] = pd.to_datetime(df1['tpep_pickup_datetime'], format='%m/%d/%Y %I:%M:%S %p')
df1['tpep_dropoff_datetime'] = pd.to_datetime(df1['tpep_dropoff_datetime'], format='%m/%d/%Y %I:%M:%S %p')
df1['tpep_pickup_datetime']
```

```
Out[13]: 0      2017-03-25 08:55:43
         1      2017-04-11 14:53:28
         2      2017-12-15 07:26:56
         3      2017-05-07 13:17:59
         5      2017-03-25 20:34:11
         ..
        22692    2017-07-16 03:22:51
        22693    2017-08-10 22:20:04
        22695    2017-08-06 16:43:59
        22697    2017-07-15 12:56:30
        22698    2017-03-02 13:02:49
Name: tpep_pickup_datetime, Length: 15265, dtype: datetime64[ns]
```

```
In [14]: # Creating a day column with the day of the week when passengers were picked up
df1['day'] = df1['tpep_pickup_datetime'].dt.day_name().str.lower()
df1["day"]
```

```
Out[14]: 0      saturday
1      tuesday
2      friday
3      sunday
5      saturday
...
22692    sunday
22693    thursday
22695    sunday
22697    saturday
22698    thursday
Name: day, Length: 15265, dtype: object
```

Defining four bins representing time of day: am_rush = [06:00–10:00]

daytime = [10:00–16:00)

pm_rush = [16:00–20:00)

nighttime = [20:00–06:00)

Creating four columns, containing binary values indicating whether a trip began during the time

```
In [15]: df1["am_rush"] = df1['tpep_pickup_datetime'].dt.hour
df1["daytime"] = df1['tpep_pickup_datetime'].dt.hour
df1["pm_rush"] = df1['tpep_pickup_datetime'].dt.hour
df1["nighttime"] = df1['tpep_pickup_datetime'].dt.hour
```

```
In [16]: def am_rush(hour):
    if 6 <= hour['am_rush'] < 10:
        val = 1
    else:
        val = 0
    return val

df1['am_rush'] = df1.apply(am_rush, axis=1)
```

```
In [17]: def daytime(hour):
    if 10 <= hour['daytime'] < 16:
        val = 1
    else:
        val = 0
    return val
df1['daytime'] = df1.apply(daytime, axis=1)
```

```
In [18]: def pm_rush(hour):
    if 16 <= hour['pm_rush'] < 20:
        val = 1
    else:
        val = 0
    return val
df1['pm_rush'] = df1.apply(pm_rush, axis=1)
```

```
In [19]: def nighttime(hour):
    if 20 <= hour['nighttime'] < 24:
        val = 1
    elif 0 <= hour['nighttime'] < 6:
        val = 1
    else:
        val = 0
    return val
df1['nighttime'] = df1.apply(nighttime, axis=1)
df1['nighttime'].head(5)
```

```
Out[19]: 0    0  
1    0  
2    0  
3    0  
5    1  
Name: nighttime, dtype: int64
```

```
In [20]: # Creating a month column with abbreviated name of the month when passengers were picked up  
df1['month'] = df1['tpep_pickup_datetime'].dt.strftime('%b').str.lower()
```

```
In [21]: # Comparing to the originating df0, 8 columns were added: tip_percent, generous, day, am_rush, daytime,  
df1.head()
```

```
Out[21]:   Unnamed:  
          0  VendorID  tpep_pickup_datetime  tpep_dropoff_datetime  passenger_count  trip_distance  Ratecode  
0    24870114        2  2017-03-25 08:55:43  2017-03-25 09:09:47            6       3.34  
1    35634249        1  2017-04-11 14:53:28  2017-04-11 15:19:58            1       1.80  
2    106203690       1  2017-12-15 07:26:56  2017-12-15 07:34:08            1       1.00  
3    38942136        2  2017-05-07 13:17:59  2017-05-07 13:48:14            1       3.70  
5    23345809        2  2017-03-25 20:34:11  2017-03-25 20:42:11            6       2.30
```

Irrelevant Columns Removal

```
In [22]: # The following columns dropped are either irrelevant, redundant, or won't be available when the model  
drop_cols = ['Unnamed: 0', 'tpep_pickup_datetime', 'tpep_dropoff_datetime',  
             'payment_type', 'trip_distance', 'store_and_fwd_flag', 'payment_type',  
             'fare_amount', 'extra', 'mta_tax', 'tip_amount', 'tolls_amount',  
             'improvement_surcharge', 'total_amount', 'tip_percent']  
  
df1 = df1.drop(drop_cols, axis=1)  
df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Index: 15265 entries, 0 to 22698  
Data columns (total 15 columns):  
 #   Column           Non-Null Count  Dtype     
---  --  
 0   VendorID         15265 non-null   int64    
 1   passenger_count  15265 non-null   int64    
 2   RatecodeID       15265 non-null   int64    
 3   PULocationID    15265 non-null   int64    
 4   DOLocationID    15265 non-null   int64    
 5   mean_duration    15265 non-null   float64  
 6   mean_distance    15265 non-null   float64  
 7   predicted_fare   15265 non-null   float64  
 8   generous         15265 non-null   int64    
 9   day               15265 non-null   object   
 10  am_rush          15265 non-null   int64    
 11  daytime          15265 non-null   int64    
 12  pm_rush          15265 non-null   int64    
 13  nighttime         15265 non-null   int64    
 14  month             15265 non-null   object   
dtypes: float64(3), int64(10), object(2)  
memory usage: 1.9+ MB
```

Variable Encoding

Converting numerical columns that contain categorical information to string, then to binary using one-hot encoding.

```
In [23]: cols_to_str = ['RatecodeID', 'PULocationID', 'DOLocationID', 'VendorID']
for col in cols_to_str:
    df1[col] = df1[col].astype('str')
# Converting categoricals to binary/dummies
df2 = pd.get_dummies(df1, drop_first=True)
df2
```

```
Out[23]:
```

	passenger_count	mean_duration	mean_distance	predicted_fare	generous	am_rush	daytime	pm_rush
0	6	22.847222	3.521667	16.434245	1	1	0	0
1	1	24.470370	3.108889	16.052218	1	0	1	0
2	1	7.250000	0.881429	7.053706	0	1	0	0
3	1	30.250000	3.700000	18.731650	1	0	1	0
5	6	11.855376	2.052258	10.441351	1	0	0	0
...
22692	1	18.016667	5.700000	19.426247	0	0	0	0
22693	1	8.095370	1.062778	7.300146	1	0	0	0
22695	1	59.560417	18.757500	52.000000	1	0	0	1
22697	1	16.650000	2.077500	11.707049	0	0	1	0
22698	1	9.405556	1.476970	8.600969	0	0	1	0

15265 rows × 347 columns

There are only a few hundred columns here, scalability concern is low. We will not need to drop less frequent categories or perform hashing.

Evaluation metric

```
In [24]: # Evaluate class balance of 'generous' col
df2['generous'].value_counts(normalize=True)
```

```
Out[24]:
```

generous	
0	0.511169
1	0.488831
Name:	proportion, dtype: float64

The dataset is nearly balanced.

Ethics

- False positives (the model predicts a tip >= 20%, but the customer gives less)
- False negatives (the model predicts a tip < 20%, but the customer gives more)

If some forms of the prediction of this model is visible to the drivers: False positives are worse for cab drivers, because they would pick up a customer expecting a good tip and then not receive one. False negatives are worse for customers, because a customer might not be picked up even though he/she would tip generously

Our metric, which is F_1 score, should place equal weight on true postives and false positives, and so therefore on precision and recall.

4. Modeling

Data Split

```
In [25]: # Isolate target variable
y = df2['generous']
# Isolate the features
X = df2.drop('generous', axis=1)

# Putting 20% of the samples into test, stratifying data, setting random state
X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y, test_size=0.2, random_state=27)
```

Random Forest Model

```
In [26]: rf = RandomForestClassifier(random_state=27, n_jobs=-1)

cv_params = {
    # Number of trees
    'n_estimators': [100, 300],
    # Tree depth - Overfitting control
    'max_depth': [None, 10, 20],
    # How many samples to consider for a split
    'min_samples_split': [2, 10],
    # Minimum samples per leaf - 1 = pure, higher = smoother
    'min_samples_leaf': [1, 4],
    # Number of features trying at each split
    'max_features': ['sqrt', 0.3],
    # Subsample rows for each tree
    'max_samples': [None, 0.7],
    # Potential class imbalance
    'class_weight': [None, 'balanced']
}

scoring = ['accuracy', 'precision', 'recall', 'f1']

rf_grid = GridSearchCV(
    estimator=rf,
    param_grid=cv_params,
    scoring=scoring,
    cv=4,
    refit='f1',
    n_jobs=-1
)
```

```
In [27]: %%time
rf_grid.fit(X_train, y_train)
```

CPU times: total: 6.27 s
Wall time: 6min 39s

Out[27]:

GridSearchCV		
▼ Parameters		
estimator		RandomForestC...ndom_state=27)
param_grid		{'class_weight': [None, 'balanced'], 'max_depth': [None, 10, ...], 'max_features': ['sqrt', 0.3], 'max_samples': [None, 0.7], ...}
scoring		['accuracy', 'precision', ...]
n_jobs		-1
refit		'f1'
cv		4
verbose		0
pre_dispatch		'2*n_jobs'
error_score		nan
return_train_score		False

▼ best_estimator_: RandomForestClassifier

RandomForestClassifier(max_depth=10, max_features=0.3, max_samples=0.7,
min_samples_leaf=4, n_jobs=-1, random_state=27)

▼ RandomForestClassifier		
▼ Parameters		
n_estimators		100
criterion		'gini'
max_depth		10
min_samples_split		2
min_samples_leaf		4
min_weight_fraction_leaf		0.0
max_features		0.3
max_leaf_nodes		None
min_impurity_decrease		0.0
bootstrap		True
oob_score		False
n_jobs		-1
random_state		27
verbose		0
warm_start		False
class_weight		None

ccp_alpha	0.0
max_samples	0.7
monotonic_cst	None

```
In [28]: rf_grid.best_score_
```

```
Out[28]: np.float64(0.7365020307243173)
```

```
In [29]: rf_grid.best_params_
```

```
Out[29]: {'class_weight': None,
'max_depth': 10,
'max_features': 0.3,
'max_samples': 0.7,
'min_samples_leaf': 4,
'min_samples_split': 2,
'n_estimators': 100}
```

```
In [30]: # Function to oupt all scores of the model
```

```
def make_results(model_name:str, model_object, metric:str):
    # Dictionary that maps input metric to actual metric name in GridSearchCV
    metric_dict = {'precision': 'mean_test_precision',
                  'recall': 'mean_test_recall',
                  'f1': 'mean_test_f1',
                  'accuracy': 'mean_test_accuracy',
                  }

    cv_results = pd.DataFrame(model_object.cv_results_)

    # Isolate the row of the df with the max(metric) score
    best_estimator_results = cv_results.iloc[cv_results[metric_dict[metric]].idxmax(), :]

    # Extract Accuracy, precision, recall, and f1 score from that row
    f1 = best_estimator_results.mean_test_f1
    recall = best_estimator_results.mean_test_recall
    precision = best_estimator_results.mean_test_precision
    accuracy = best_estimator_results.mean_test_accuracy

    table = pd.DataFrame({'model': [model_name],
                          'precision': [precision],
                          'recall': [recall],
                          'F1': [f1],
                          'accuracy': [accuracy],
                          },
                          )

    return table
```

```
In [31]: results = make_results('RF CV', rf_grid, 'f1')
results
```

	model	precision	recall	F1	accuracy
0	RF CV	0.693205	0.785593	0.736502	0.725188

The Cross Validation shows some valid prediction power. But I will train a different model and pick the best one.

```
In [32]: def get_test_scores(model_name:str, preds, y_test_data):  
  
    accuracy = accuracy_score(y_test_data, preds)  
    precision = precision_score(y_test_data, preds)  
    recall = recall_score(y_test_data, preds)  
    f1 = f1_score(y_test_data, preds)  
  
    table = pd.DataFrame({'model': [model_name],  
                          'precision': [precision],  
                          'recall': [recall],  
                          'F1': [f1],  
                          'accuracy': [accuracy]})  
  
    return table
```

```
In [33]: rf_preds = rf_grid.best_estimator_.predict(X_test)  
rf_test_scores = get_test_scores('RF test', rf_preds, y_test)  
results = pd.concat([results, rf_test_scores], axis=0)  
results
```

```
Out[33]:   model  precision      recall        F1  accuracy  
0  RF CV    0.693205  0.785593  0.736502  0.725188  
0  RF test   0.698460  0.790214  0.741509  0.730757
```

Test results are slightly higher comparing to CV results. The Random Forest model achieves an F1 score of 0.74 on the test set, indicating acceptable performance given the inherent noise and unpredictability of tipping behavior. In production systems, it is difficult for behavioral predictions to exceed 0.80 F1 due to missing contextual factors not present in the dataset such as customer mood, driver-customer interactions. Therefore, such a performance could be considered useful if it meaningfully improves business outcomes.

XGBoost Model

```
In [34]: xgb = XGBClassifier(objective='binary:logistic', random_state=0, n_jobs=-1, eval_metric='logloss', tree_  
  
#This grid is a compromise from a larger grid, as my local machine does not have enough computing power  
cv_params = {  
    'learning_rate': [0.05, 0.1],  
    # Tree depth: shallow/medium/deeper  
    'max_depth': [3, 5, 7],  
    # Overfitting control  
    'min_child_weight': [1, 3],  
    # Number of trees  
    'n_estimators': [300, 500],  
    # Row subsampling  
    'subsample': [0.8, 1.0],  
    # Column subsampling per tree  
    'colsample_bytree': [0.8, 1.0],  
    # Regularization knobs  
    'gamma': [0, 0.5],  
    'reg_lambda': [1.0],      #L2  
    'reg_alpha': [0, 0.01]    #L1  
}
```

```
scoring = ['accuracy', 'precision', 'recall', 'f1']

xgb_grid = GridSearchCV(
    estimator=xgb,
    param_grid=cv_params,
    scoring=scoring,
    cv=4,
    refit='f1',
    n_jobs=-1,
    verbose=1
)
```

In [35]: `%%time`
`xgb_grid.fit(X_train, y_train)`

```
Fitting 4 folds for each of 384 candidates, totalling 1536 fits
CPU times: total: 15.1 s
Wall time: 6min 39s
```

Out[35]:

GridSearchCV		
▼ Parameters		
estimator	XGBClassifier(base_score=None, booster=None, callbacks=None, colsample_bytree=1.0, device=None, early_stopping_rounds=None, enable_categorical=False, eval_metric='logloss', feature_types=None, feature_weights=None, gamma=0.5, grow_policy=None, importance_type=None, interaction_constraints=None, learning_rate=0.05, max_bin=None, max_cat_threshold=None, max_cat_to_onehot=None, max_delta_step=None, max_depth=3, max_leaves=None, min_child_weight=3, missing=nan, monotone_constraints=None, multi_strategy=None, n_estimators=300, n_jobs=-1, num_parallel_tree=None, ...)	
param_grid	{'colsample_bytree': [0.8, 1.0], 'gamma': [0, 0.5], 'learning_rate': [0.05, 0.1], 'max_depth': [3, 5, ...], ...}	
scoring	['accuracy', 'precision', ...]	
n_jobs	-1	
refit	'f1'	
cv	4	
verbose	1	
pre_dispatch	'2*n_jobs'	
error_score	nan	
return_train_score	False	
▼ best_estimator_: XGBClassifier		
XGBClassifier(base_score=None, booster=None, callbacks=None, colsample_bylevel=None, colsample_bynode=None, colsample_bytree=1.0, device=None, early_stopping_rounds=None, enable_categorical=False, eval_metric='logloss', feature_types=None, feature_weights=None, gamma=0.5, grow_policy=None, importance_type=None, interaction_constraints=None, learning_rate=0.05, max_bin=None, max_cat_threshold=None, max_cat_to_onehot=None, max_delta_step=None, max_depth=3, max_leaves=None, min_child_weight=3, missing=nan, monotone_constraints=None, multi_strategy=None, n_estimators=300, n_jobs=-1, num_parallel_tree=None, ...)		
▼ XGBClassifier		
▼ Parameters		
objective	'binary:logistic'	
base_score	None	
booster	None	
callbacks	None	
colsample_bylevel	None	
colsample_bynode	None	
colsample_bytree	1.0	
device	None	

early_stopping_rounds		None
enable_categorical		False
eval_metric		'logloss'
feature_types		None
feature_weights		None
gamma		0.5
grow_policy		None
importance_type		None
interaction_constraints		None
learning_rate		0.05
max_bin		None
max_cat_threshold		None
max_cat_to_onehot		None
max_delta_step		None
max_depth		3
max_leaves		None
min_child_weight		3
missing		nan
monotone_constraints		None
multi_strategy		None
n_estimators		300
n_jobs		-1
num_parallel_tree		None
random_state		0
reg_alpha		0
reg_lambda		1.0
sampling_method		None
scale_pos_weight		None
subsample		0.8
tree_method		'hist'
validate_parameters		None
verbosity		None

In [36]: `xgb_grid.best_score_`

```
Out[36]: np.float64(0.7343495003401214)
```

```
In [37]: xgb_grid.best_params_
```

```
Out[37]: {'colsample_bytree': 1.0,
          'gamma': 0.5,
          'learning_rate': 0.05,
          'max_depth': 3,
          'min_child_weight': 3,
          'n_estimators': 300,
          'reg_alpha': 0,
          'reg_lambda': 1.0,
          'subsample': 0.8}
```

```
In [39]: xgb_cv_results = make_results('XGB CV', xgb_grid, 'f1')
      results = pd.concat([results, xgb_cv_results], axis=0)
      results
```

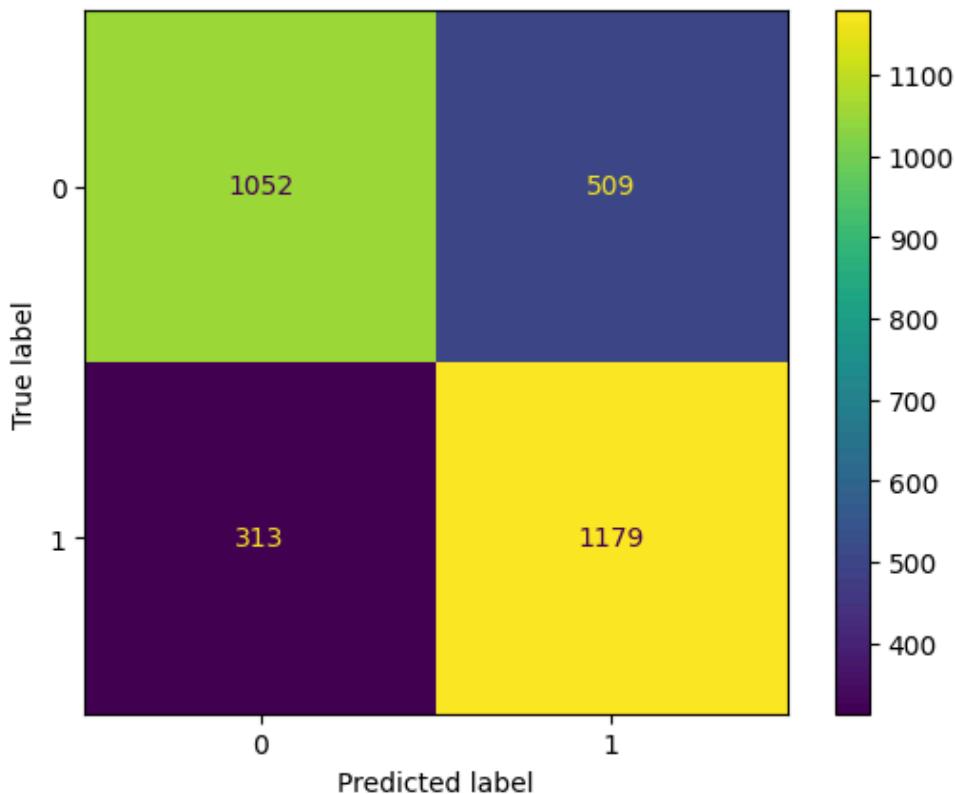
Out[39]:	model	precision	recall	F1	accuracy
0	RF CV	0.693205	0.785593	0.736502	0.725188
0	RF test	0.698460	0.790214	0.741509	0.730757
0	XGB CV	0.692520	0.781574	0.734350	0.723551

```
In [41]: xgb_preds = xgb_grid.best_estimator_.predict(X_test)
xgb_test_scores = get_test_scores('XGB test', xgb_preds, y_test)
results = pd.concat([results, xgb_test_scores], axis=0)
results
```

Out[41]:	model	precision	recall	F1	accuracy
0	RF CV	0.693205	0.785593	0.736502	0.725188
0	RF test	0.698460	0.790214	0.741509	0.730757
0	XGB CV	0.692520	0.781574	0.734350	0.723551
0	XGB test	0.698516	0.788874	0.740951	0.730429

Since we chose F1 score as the optimization metric, the two models have identical performances, with Random Forest being the slightly better one. XGBoost is a high-capacity model that requires careful hyperparameter tuning to avoid overfitting noisy signals. However, the dataset contains many sparse dummy variables and significant behavioral noise. Also, XGBoost is calculated mostly sequentially, making a larger hyperparameter search space impractical under computational constraints. Therefore, XGBoost was unable to reach its optimal configuration. In contrast, Random Forest is more robust to noise and performs well even with minimal tuning, which explains why it outperformed XGBoost in this case.

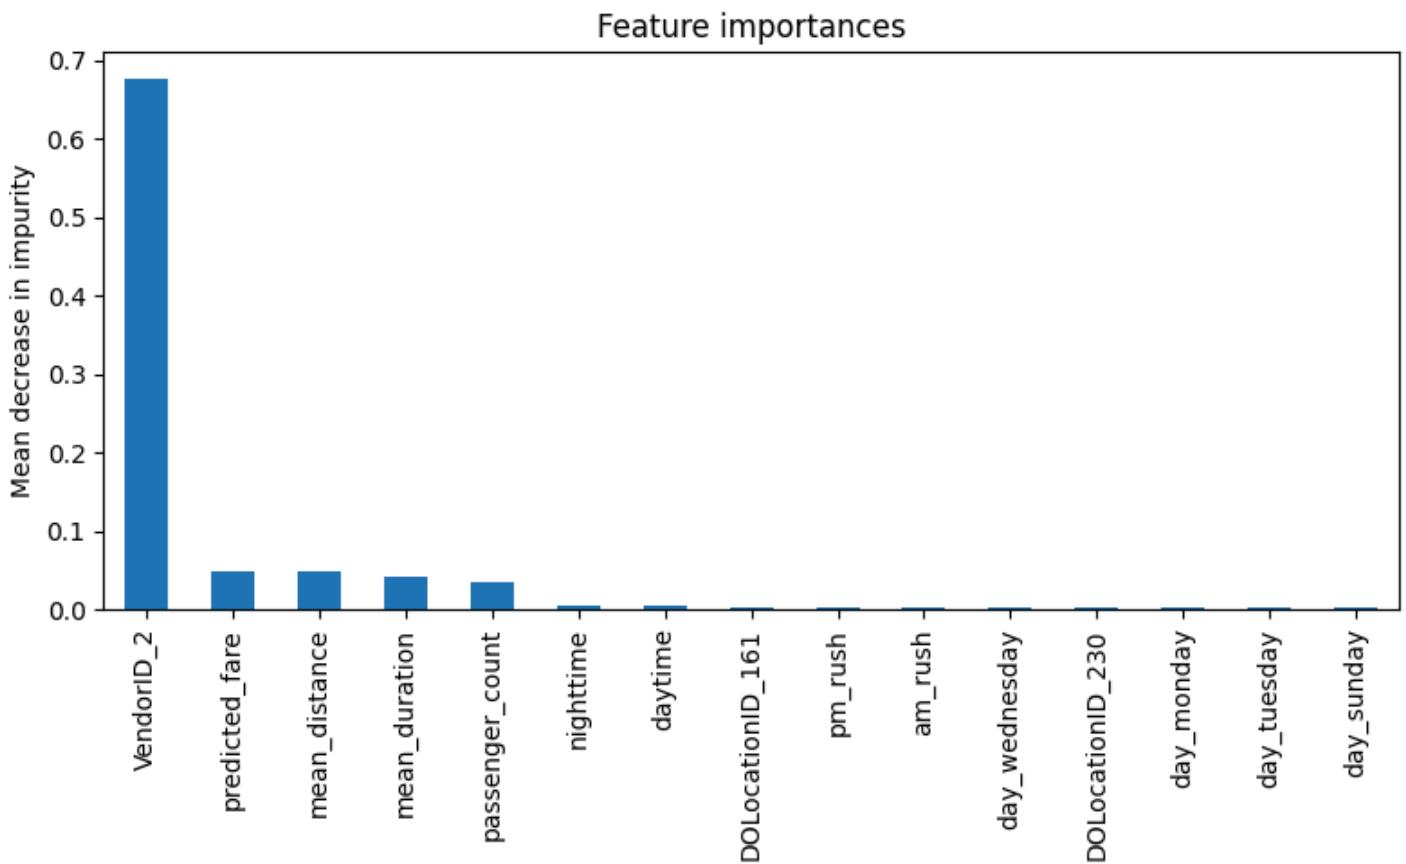
Random Forest Confusion Matrix



As the champion model, it is much more likely to predict a false positive than false negative. In another word,, Type I errors are more common. In production this could be less desirable. In more cases drivers would be disappointed by a low tip, expecting a generous one, rather than be pleasantly surprised by a generous tip.

```
In [48]: importances = rf_grid.best_estimator_.feature_importances_
rf_importances = pd.Series(importances, index=X_test.columns)
rf_importances = rf_importances.sort_values(ascending=False)[:15]

fig, ax = plt.subplots(figsize=(8,5))
rf_importances.plot.bar(ax=ax)
ax.set_title('Feature importances')
ax.set_ylabel('Mean decrease in impurity')
fig.tight_layout();
```



We know that **VendorID**, **predicted_fare**, **mean_distance**, **mean_duration**, and **passenger count** are the most important features, but we are not sure how they influence tipping. Random forest is not the most transparent machine learning algorithm. Further statistical tests on features such as **VendorID** might reveal more insights.

4. Conclusion

Overall it is a good model. It correctly predicts about 73% of the customers, which is 46% better than random guessing. Additionally, we could add more features, for example people are more likely to round up their tip, which could affect our calculation. In the future, another model that should be prioritized is LightGBM, which is faster and handles large categorical variable better. With a large dataset and budget to cloud computing, we could also test out other machine learning models such as CatBoost and TabNet.